

CS3319 Foundations of Data Science

# 4. Locality Sensitive Hashing

Jiaxin Ding

John Hopcroft Center



上海交通大学  
约翰·霍普克罗夫特  
计算机科学中心

John Hopcroft Center for Computer Science



# Text Similarity

- Similarity check for a paper with all the published papers.

Running title • August 2012 • Vol. XXI, No. 1

---

## Article Title

JOHN SMITH\*  
University of California  
john@smith.com

**Abstract**

*Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate n, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.*

### I. INTRODUCTION

**L**orem ipsum dolor sit amet, consectetur adipiscing elit. Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula

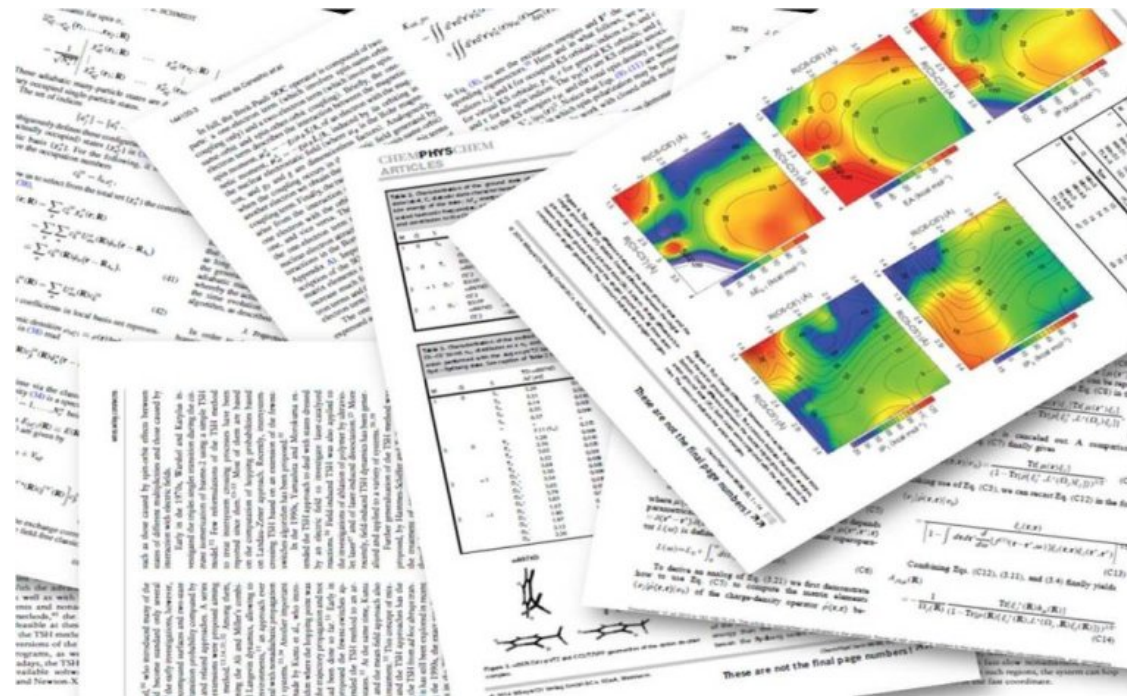
feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

### II. METHODS

Maecenas sed ultricies felis. Sed imperdiet dictum arcu a egestas.

- Donec dolor arcu, rutrum id molestie in, viverra sed diam
- Curabitur feugiat
- turpis sed auctor facilisis
- arcu eros accumsan lorem, at posuere mi diam sit amet tortor
- Fusce fermentum, mi sit amet euismod rutrum
- sem lorem molestie diam, iaculis aliquet sapien tortor non nisi
- Pellentesque bibendum pretium aliquet Quisque ullamcorper placerat ipsum. Cras nibh. Morbi vel justo vitae lacus tincidunt ultris

\*A thank you or further information



# Task: Finding Similar Documents

- **Goal:** Given a **large number** ( **$N$**  in the millions or billions) of documents, find “**near duplicate**” pairs
- Challenges:
  - How to define the **similarity**?
    - Many small pieces of one document can appear **out of order** in another.
  - How to compute **efficiently**?
    - Documents are so **large** or so many that they cannot fit in main memory
    - **Too many** documents to compare **all pairs**. E.g. 1 million documents, we have  $10^{12}$  pairs, if we compare  $10^6$  per second, it takes about 10 days.


# Problem Definition

- **Given:** High dimensional data points (e.g. Bag of Words)  $x_1, x_2, \dots$
- A **distance function**  $d(x_1, x_2)$
- **Goal:** Find all pairs of data points  $(x_i, x_j)$  that are within some distance threshold  $d(x_i, x_j) \leq s$
- Naïve solution would take  $O(N^2)$ 
  - where  $N$  is the number of data points
- This can be done in  $O(N)$ !

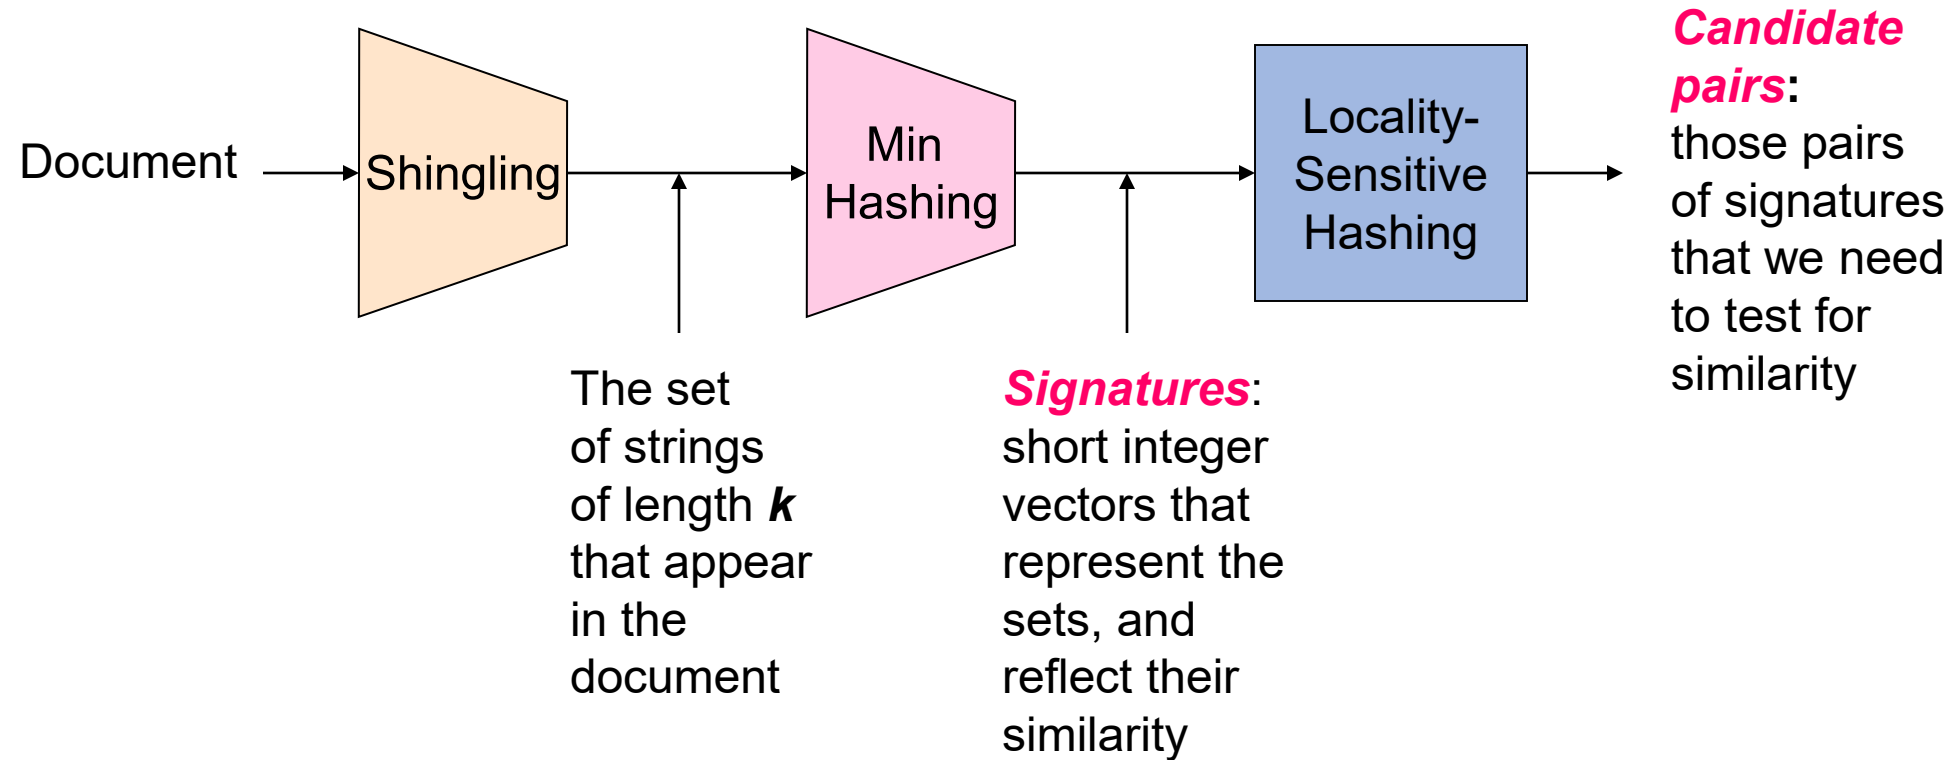
# Key Idea

- **Hashing**

- Throw items into buckets using several different **hash functions**.
- Examine only those pairs of items that **share a bucket** for at least one of these hashings.

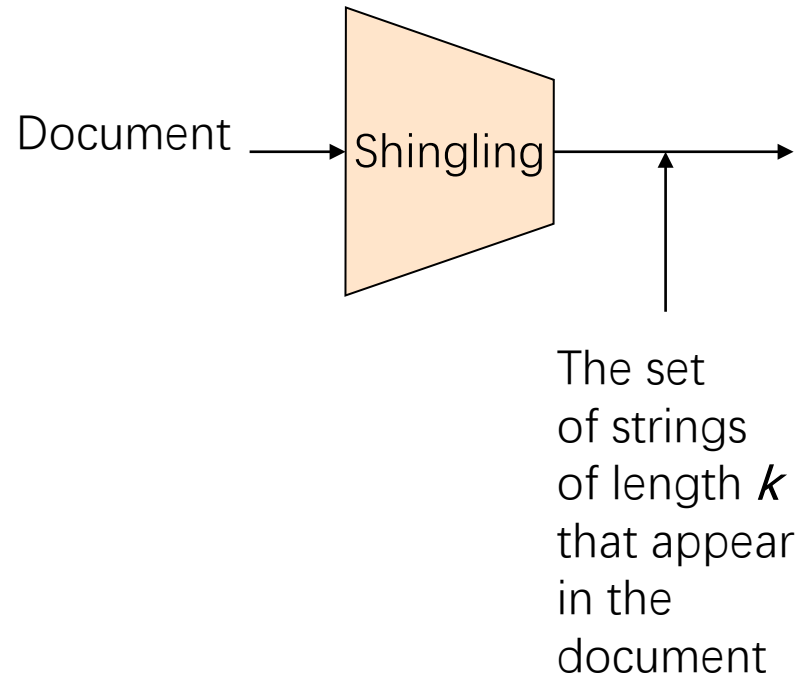


# The Big Picture: 3 Steps for Similar Documents



# Shingling

- Step 1: Shingling: Convert documents to sets



# Documents as High Dimensional Data

- Step 1: **Shingling**: Convert documents to sets
- Simple approaches:
  - Document = set of **words** appearing in document
  - Document = set of “**important**” words
- Need to account for **ordering** of words!
- A different way: Shingles!



# Define: Shingles

- A **k-shingle** (or k-gram) for a document is a sequence of k tokens that appears in the document
  - Tokens can be characters, or words, depending on the application
- Represent a document by the set of its k-shingles
- **Example:**  $k = 2$ ; document  $D_1 = abcab$   
Set of 2-shingles:  $S(D_1) = \{ab, bc, ca\}$




# Similarity Metric for Shingles

- Document  $D_1$  is a set of its k-shingles  $C_1 = S(D_1)$
- A natural similarity measure is the **Jaccard similarity**:

$$\text{sim}(D_1, D_2) = |C_1 \cap C_2| / |C_1 \cup C_2|$$

- **Jaccard distance**:

$$d(D_1, D_2) = 1 - \text{sim}(D_1, D_2)$$



$$\begin{aligned} \text{E.g. } |C_1 \cup C_2| &= 8 \\ |C_1 \cap C_2| &= 4 \\ \text{sim}(D_1, D_2) &= 0.5 \end{aligned}$$

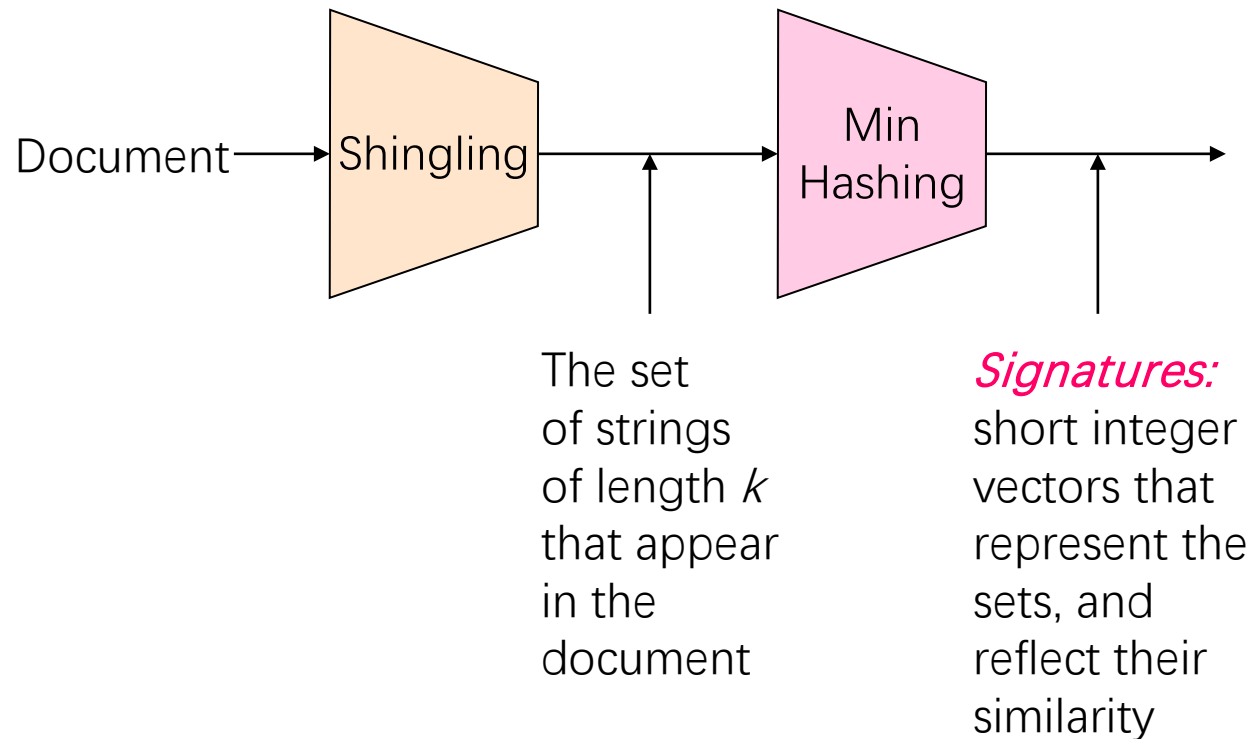
# Set Representation

- Encode sets with 0/1 vectors
- Rows = elements (shingles)
- Columns = sets (documents)
  - 1 in row  $e$  and column  $s$  if and only if  $e$  is a member of  $s$
  - Column similarity is the Jaccard similarity
  - Typical matrix is sparse!

		Documents			
		text1	text2		
Shingles	ab	1	1	1	0
	ac	1	1	0	1
		0	1	0	1
		0	0	0	1
		1	0	0	1
		1	1	1	0
		1	0	1	0

# MinHashing

- Step 2: Min-hashing: Convert large sets to short signatures, while preserving similarity



# Signatures

- Key idea: “hash” each column  $C$  to a small **signature**  $h(C)$ , such that:
  - (1)  $h(C)$  is **small** enough
  - (2)  $\text{sim}(C1, C2)$  is the same as the “**similarity**” of signatures  $h(C1)$  and  $h(C2)$

- **Goal:** Find a hash function  $h(\cdot)$  such that:
  - If  $\text{sim}(C1, C2)$  is high, then with high prob.  $h(C1) = h(C2)$
  - If  $\text{sim}(C1, C2)$  is low, then with high prob.  $h(C1) \neq h(C2)$

- Hash function for the Jaccard similarity: **Min-Hashing**

# Min-Hashing

- Imagine the rows of the boolean matrix permuted under **random permutation**  $\pi$

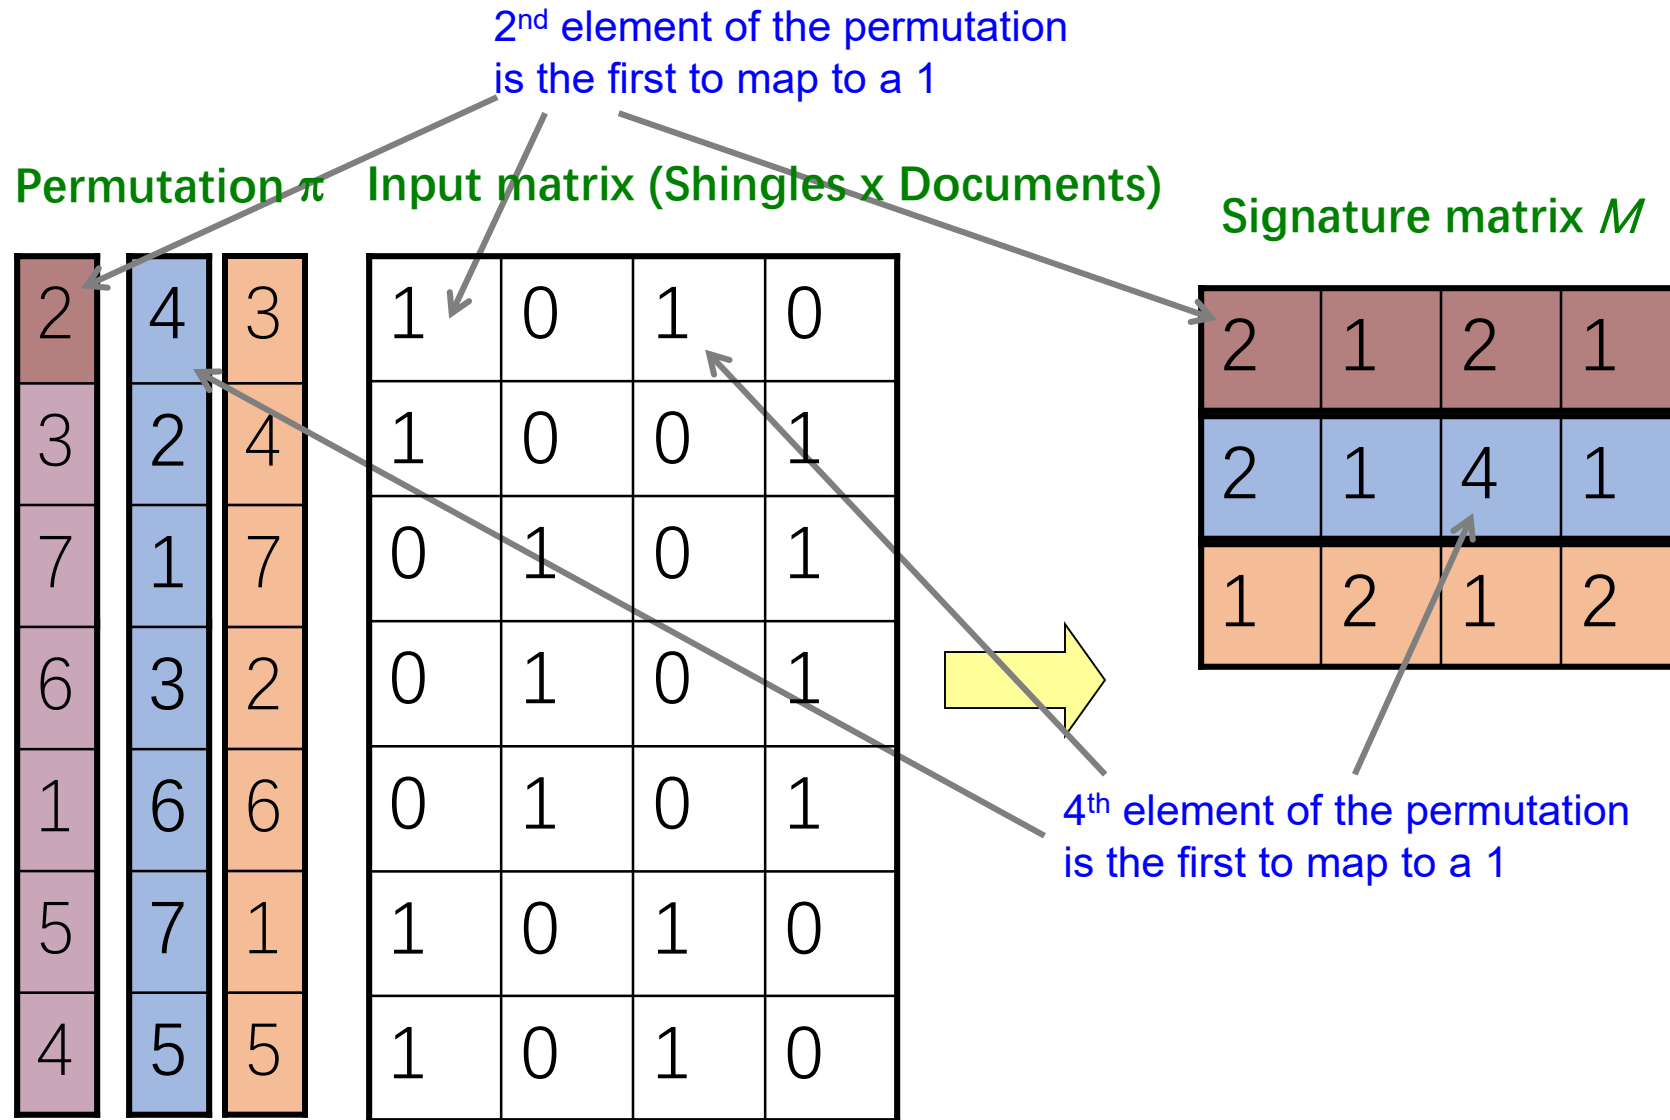
2	1	0	1	0
3	1	0	0	1
7	0	1	0	1
6	0	1	0	1
1	0	1	0	1
5	1	0	1	0
4	1	0	1	0

- Define minhash function  $h_{\pi}(C)$  = the **index** of the **first** (in the permuted order  $\pi$ ) row in which column  $C$  has value **1**:

$$h_{\pi}(C) = \min \pi(C)$$

- Use **independent hash** functions to create a **signature** of a **column**

# Min-Hashing Example



# The Min-Hash Property

- Choose a random permutation  $\pi$
- Claim:  $\Pr[h_\pi(C_1) = h_\pi(C_2)] = \text{sim}(C_1, C_2)$
- Why?
  - Let  $X$  be a doc (set of shingles),  $y \in X$  is a shingle
  - Then:  $\Pr[\pi(y) = \min(\pi(X))] = 1/|X|$ 
    - It is **equally likely** that any  $y \in X$  is mapped to the min element
  - Let  $y$  satisfy  $\pi(y) = \min(\pi(C_1 \cup C_2))$
  - Then either:  $\pi(y) = \min(\pi(C_1))$  if  $y \in C_1$ , or  $\pi(y) = \min(\pi(C_2))$  if  $y \in C_2$
  - So the prob. that **both are true** is the prob.  $y \in C_1 \cap C_2$
  - $\Pr[\pi(y) = \min(\pi(C_1)) = \min(\pi(C_2))] = |C_1 \cap C_2| / |C_1 \cup C_2| = \text{sim}(C_1, C_2)$

0	0
0	0
<b>1</b>	<b>1</b>
0	0
0	1
1	0

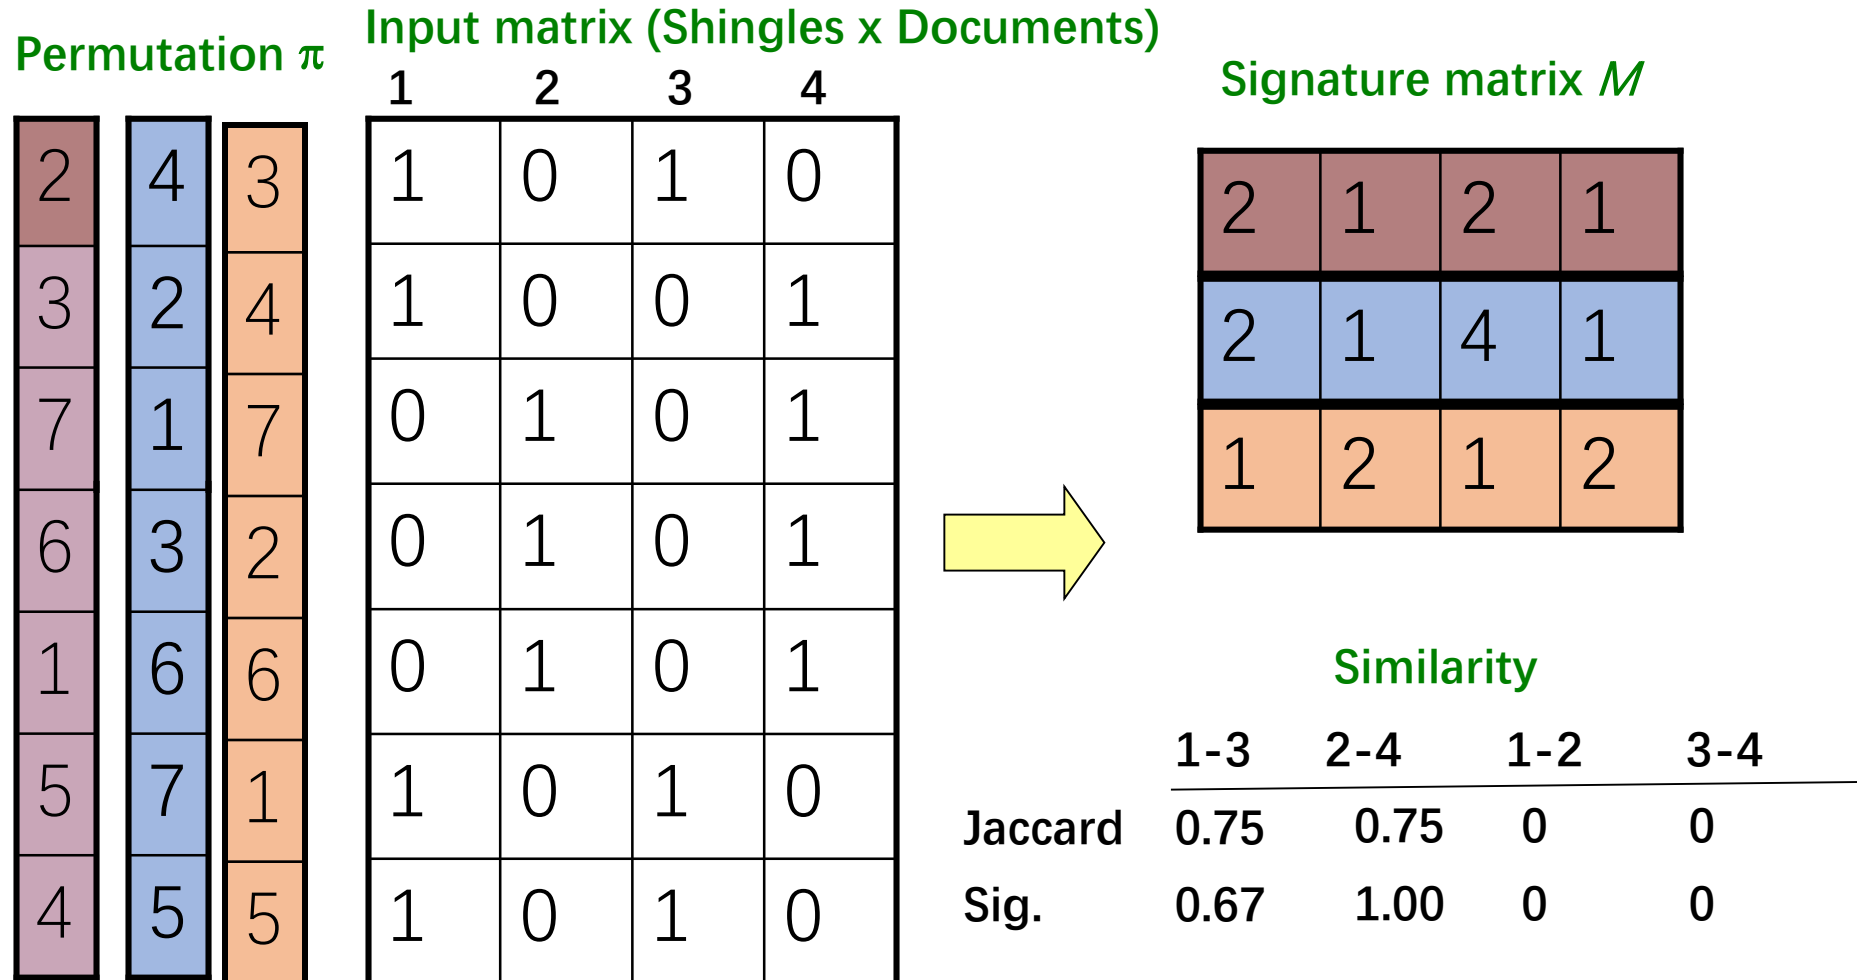
One of the two  
cols had to have  
1 at position  $y$



# Similarity for Signatures

- *We know:*  $\Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = \text{sim}(C_1, C_2)$
- Now generalize to **multiple** hash functions
- The **similarity** of two signatures is the **fraction** of the hash functions in which they agree
- **Note:** Because of the Min-Hash property, the similarity of columns is the same as the expected similarity of their signatures, with **expected error** of  $O\left(\frac{1}{\sqrt{k}}\right)$ ,  $k$  is the number of hash functions.

# Min-Hashing Example



# Implementation Trick

- **Permuting** rows is **complicated**, we only need the minimum hashing
- **Row hashing**
  - Pick  $K = 100$  hash functions  $h_i$
  - Ordering under  $h_i$  gives a random row permutation
- **One-pass implementation**
  - For each column  $C$  and hash func.  $h_i$
  - Initialize all  $M(i, C) = \infty$ , to store the **smallest** hashing value of a document under  $h_i$
  - Scan rows looking for 1s
    - Suppose row  $j$  has 1 in column  $C$
    - Then for each  $h_i$  :
      - If  $h_i(j) < M(i, C)$ , then  $M(i, C) = h_i(j)$ .

How to pick a random hash function  $h(x)$ ?

**Universal hashing:**

$h_{a,b}(x) = ((a \cdot x + b) \bmod p) \bmod N$   
where:

$a, b$  ... random integers

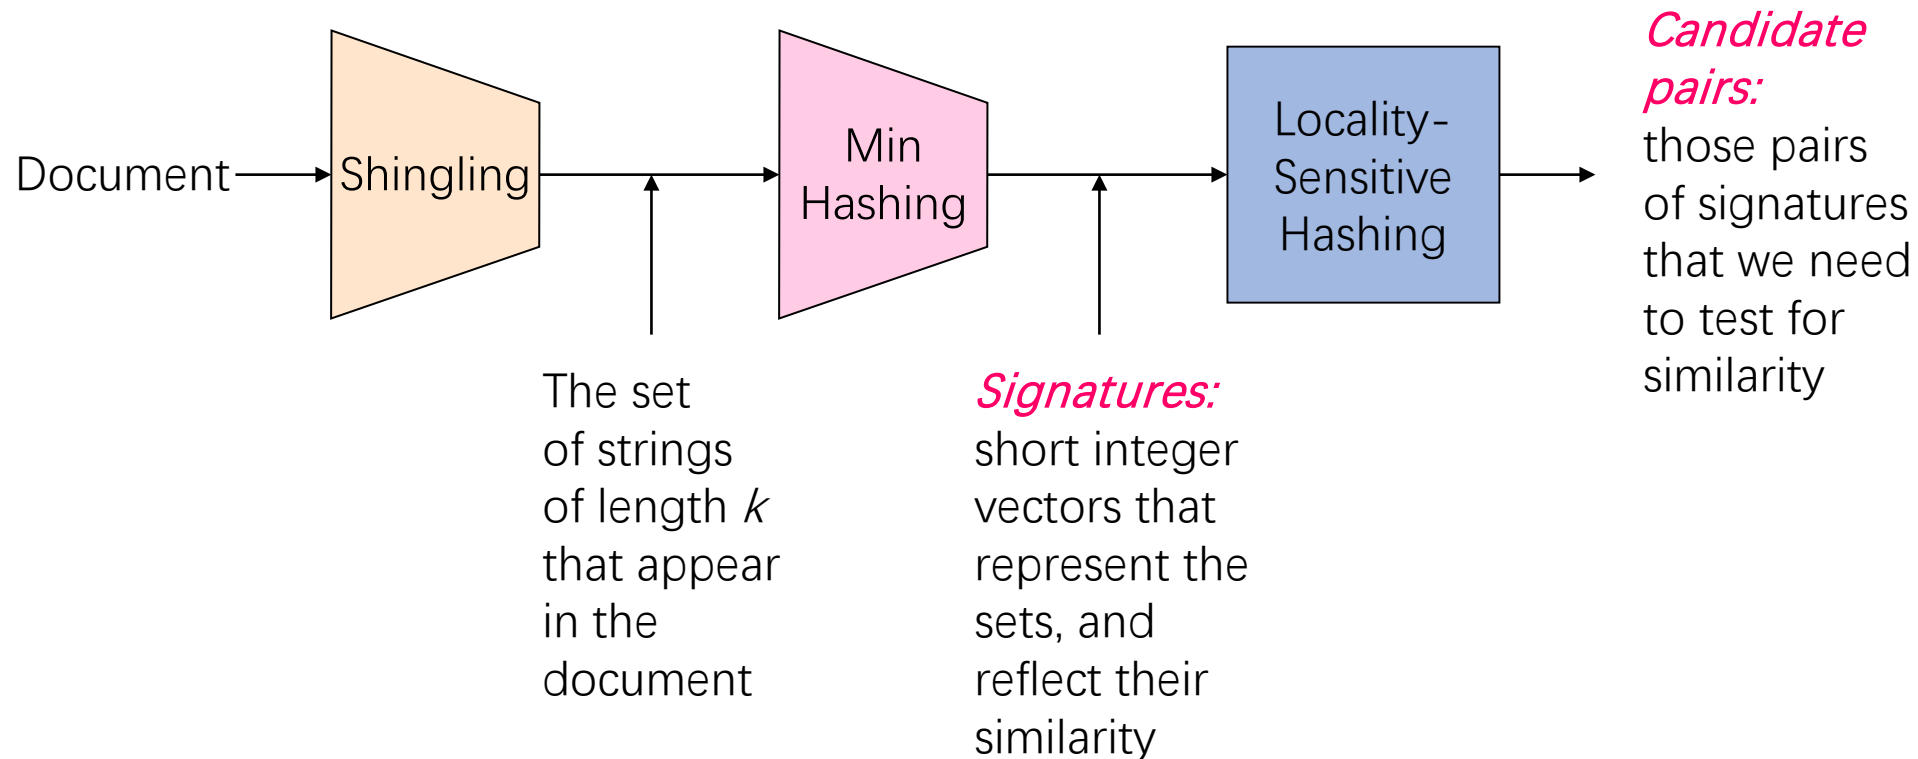
$p$  ... prime number ( $p > N$ )

Signature matrix  $M$

		$C$			
$i$		2	1	2	1
		2	1	4	1
		1	2	1	2

# Locality Sensitive Hashing


- Step 3: Locality-Sensitive Hashing:  
Focus on pairs of signatures likely to be from similar documents




# Locality Sensitive Hashing

- **Goal:** Find documents with Jaccard similarity at least  $s$  (for some similarity threshold, e.g.,  $s = 0.8$ )
- **LSH – General idea:** Use a **function**  $f(x, y)$  that tells whether  $x$  and  $y$  is a **candidate pair**
- **For Min-Hash matrices:**
  - **Hash** **columns** of **signature** matrix  $M$  to many **buckets**
  - Each pair of documents that hashes into the **same bucket** is a candidate pair

# Jaccard Similarity Hashing (1 signature)




# What We Want



What can we do with multiple minhash signatures?

# Jaccard Similarity Hashing



We have  $k$  hash functions.

- We consider  $C_1, C_2$  to be a candidate pair, only if they share **all** the  $k$  Minhash values (**AND**)
- We consider  $C_1, C_2$  to be a candidate pair, if they share at least **one** Minhash value (**OR**)

What can we do?




# LSH for Min-Hash

- **Key Idea:**


- Hash columns of signature matrix  $M$  several times

- **Combine OR and AND**



# Partition M into Bands

- Divide matrix  $M$  into  $b$  bands of  $r$  rows
- For each band, hash its portion of each column to a hash table (Buckets: as many as possible)
- Candidate column pairs are those that hash to the same bucket for  $\geq 1$  band
- Tune  $b$  and  $r$  to catch most similar pairs, but few non-similar pairs



# Example of Bands

- Assume the following case:
- Suppose 100,000 columns of  $M$  (100k docs)
- Signatures of 100 integers (rows)
- Choose  $b = 20$  bands of  $r = 5$  integers/band
- **Goal**: Find pairs of documents that are at least  $s = 0.8$  similar

$C_1, C_2$  are 80% Similar, **false negative** rate?

- Find pairs, similarity  $\geq s = 0.8$ , set  $b = 20, r = 5$
- Assume:  $\text{sim}(C_1, C_2) = 0.8$ 
  - Since  $\text{sim}(C_1, C_2) \geq s$ , we want  $C_1, C_2$  to be a **candidate pair**: we want them to hash to at least 1 common bucket (at least one band is identical)
- Probability  $C_1, C_2$  **identical** in one particular **band**:  $(0.8)^5 = 0.328$
- Probability  $C_1, C_2$  are **not similar** in all of the 20 bands:  
 $(1 - 0.328)^{20} = 0.00035$ 
  - i.e., about 1/3000 of the 80%-similar column pairs are **false negatives** (we miss them)
  - We would find 99.965% pairs of truly similar documents

# $C_1, C_2$ are 30% Similar, false positive rate?

- Find pairs of similarity  $\geq s=0.8$ , set  $b = 20, r = 5$
- Assume:  $\text{sim}(C_1, C_2) = 0.3$ 
  - Since  $\text{sim}(C_1, C_2) < s$  we want  $C_1, C_2$  to hash to NO common buckets (all bands should be different)
- Probability  $C_1, C_2$  identical in one particular band:  $(0.3)^5 = 0.00243$
- Probability  $C_1, C_2$  identical in at least 1 of 20 bands:
$$1 - (1 - 0.00243)^{20} = 0.0474$$
  - In other words, approximately 4.74% pairs of docs with 30% similarity end up becoming candidate pairs (false positive)
    - They are false positives since we will have to examine them (they are candidate pairs) but then it will turn out their similarity is below threshold  $s$

# b bands, r rows/band


Pick:

The number of **Min-Hashes** (rows of M)

The number of **bands** b, and


The number of **rows** r per band  
to balance **false positives/negatives**

- Columns  $C_1$  and  $C_2$  have similarity  $t$
- Pick any band (r rows)
  - Prob. that all rows in band equal  
 $= t^r$
  - Prob. that some row in band unequal  
 $= 1 - t^r$
- Prob. that no band identical  
 $= (1 - t^r)^b$
- Prob. that at least 1 band identical  
 $= 1 - (1 - t^r)^b$



# Picking r and b: The S-curve

- Picking r and b to get the best S-curve
  - <https://www.desmos.com/calculator/lzzvfjiujn?lang=zh-CN>
  - r: hashed into the same bucket, b: identified as similar.



# LSH Summary

- Tune  $M$ ,  $b$ ,  $r$  to **get** almost all **pairs** with **similar signatures**, and **eliminate** most pairs that do **not** have similar signatures
- Check in main memory that candidate pairs really do have **similar signatures**
- Optional: In another pass through data, **check** that the remaining **candidate pairs** really represent similar documents