

CS 3319 Foundations of Data Science

# 8. Streaming Algorithms

Jiaxin Ding

John Hopcroft Center



上海交通大学  
约翰·霍普克罗夫特  
计算机科学中心

John Hopcroft Center for Computer Science



# Data Streams

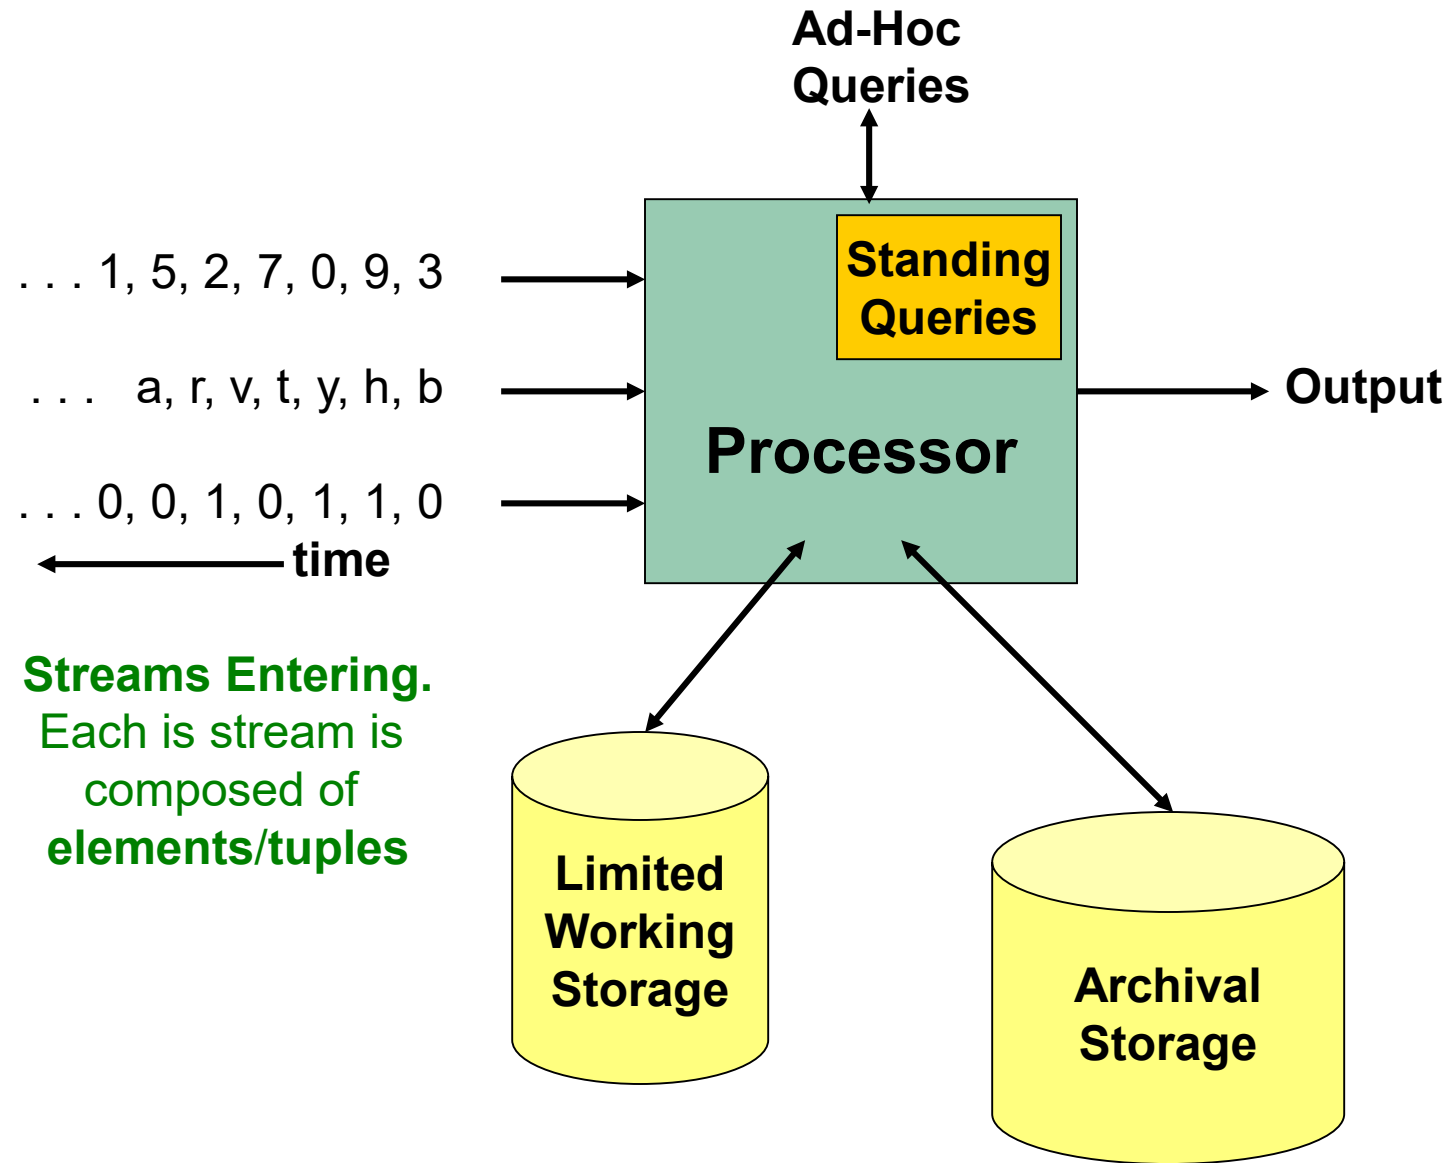
- A data stream is a **sequence** of signals used to transmit or receive information that is in the process of being **transmitted**.
  - **Infinite**
  - **Non-stationary**
- Stream Model
  - The system **cannot** store the **entire** stream
  - Input **elements(tuples)** enter at a **rapid** rate, at one or more input ports
  - Make critical calculations about the stream using a **limited** amount of memory

... 1, 5, 2, 7, 0, 9, 3

... a, r, v, t, y, h, b

... 0, 0, 1, 0, 1, 1, 0  
time

# General Stream Processing Model




It is better to use a crude approximation and know the truth, plus or minus 10 percent, than demand an exact solution and know nothing at all.

—Arthur Bloch, The Complete Murphy's Law


# Applications: Networks

- **Mining network streams**

- Finding **abnormal** patterns in sensor reading streams
- Filtering out **spam** calls in phone call streams
- Detect **denial-of-service** attacks in IP packet streams



# Applications: Internet



- **Mining query streams**

- Google wants to know what **queries** are more **frequent** today than yesterday

- **Mining click streams**

- Bytedance wants to know which of its pages are getting an **unusual** number of **hits** in the past hour

- **Mining social network news feeds**

- E.g., look for **trending topics** on Weibo

# Problems on Data Streams

- Types of queries one wants on answer on a data stream (element):
  - **Sampling data from a stream**
    - Construct a random sample
  - **Filtering a data stream**
    - Select elements with property  $x$  from the stream

# Problems on Data Streams

- Types of queries one wants on answer on a data stream (statistics):
  - **Queries over sliding windows**
    - Number of items of **type  $x$**  in the last  $k$  elements of the stream
  - **Counting distinct elements**
    - Number of **distinct** elements in the last  $k$  elements of the stream
  - **Estimating moments**
    - Estimate **avg./std. dev.** of last  $k$  elements
  - **Finding frequent elements**
    - Estimate the most **frequent** elements of the last  $k$  elements

# Sampling from a Data Stream:

## Sampling a fixed-size sample



# Maintaining a fixed-size sample

- Suppose we need to maintain a **random sample**  $S$  of size exactly  $s$  tuples
  - E.g., main memory size constraint
- Suppose at time  $n$  we have seen  $n$  items
  - Each item is in the sample  $S$  with **equal prob.**  $s/n$

**How to think about the problem: say  $s = 2$**

**Stream:** a x c y z | k c d e g...

At  $n = 5$ , each of the first 5 tuples is included in the sample  $S$  with equal prob.

At  $n = 7$ , each of the first 7 tuples is included in the sample  $S$  with equal prob.

**Q: How to achieve?**

# Solution: Fixed Size Sample

- **Algorithm**

Store all the first  $s$  elements of the stream to  $\mathcal{S}$

- Suppose we have seen  $n-1$  elements, and now the  $n^{\text{th}}$  element arrives ( $n > s$ )
  - With probability  $s/n$ , keep the  $n^{\text{th}}$  element, else discard it
  - If we picked the  $n^{\text{th}}$  element, then it replaces one of the  $s$  elements in the sample  $\mathcal{S}$ , picked uniformly at random

- This algorithm maintains a sample  $\mathcal{S}$  with the desired property:
  - After  $n$  elements, the sample contains each element seen so far with probability  $s/n$

# Proof: By Induction

- **We prove this by induction:**

- Assume that after  $n$  elements, the sample contains each element seen so far with probability  $s/n$
- We need to show that after seeing element  $n+1$  the sample maintains the property
  - Sample contains each element seen so far with probability  $s/(n+1)$

- **Base case:**

- After we see  $n=s$  elements the sample **S** has the desired property
  - Each out of  $n=s$  elements is in the sample with probability  $s/s = 1$

# Proof: By Induction

- **Inductive hypothesis:** After  $n$  elements, the sample  $\mathcal{S}$  contains each element seen so far with prob.  $s/n$
- Now element  $n+1$  arrives

- **Inductive step:** For elements already in  $\mathcal{S}$ , probability that the algorithm keeps it in  $\mathcal{S}$  is:

$$\underbrace{\left(1 - \frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ discarded}} + \underbrace{\left(\frac{s}{n+1}\right)}_{\text{Element } n+1 \text{ not discarded}} \underbrace{\left(\frac{s-1}{s}\right)}_{\text{Element in the sample not picked}} = \frac{n}{n+1}$$

- Time  $n \rightarrow n+1$ , tuple stayed in  $\mathcal{S}$  with prob.  $n/(n+1)$
- At time  $n$ , tuples in  $\mathcal{S}$  were there with prob.  $s/n$
- So prob. tuple is in  $\mathcal{S}$  at time  $n+1 = \frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

# Filtering Data Streams

# Applications

- Email **spam filtering**
  - We know 1 billion “**good**” email addresses
  - If an email comes from one of these, it is **NOT** spam
- **Publish-subscribe** systems
  - You are collecting lots of messages
  - People express interest in certain sets of **keywords**
  - Determine whether each message matches user’s **interest**

# Filtering Data Streams

- Each element of data stream is a tuple
- Given a list of keys  $S=[key_1, key_2, \dots]$
- **Determine which tuples of stream are in  $S$**
- Obvious solution: store and compare
  - But suppose we **do not have enough memory** to store all of  $S$
  - The **complexity** is  $O(S)$ , which can be big.

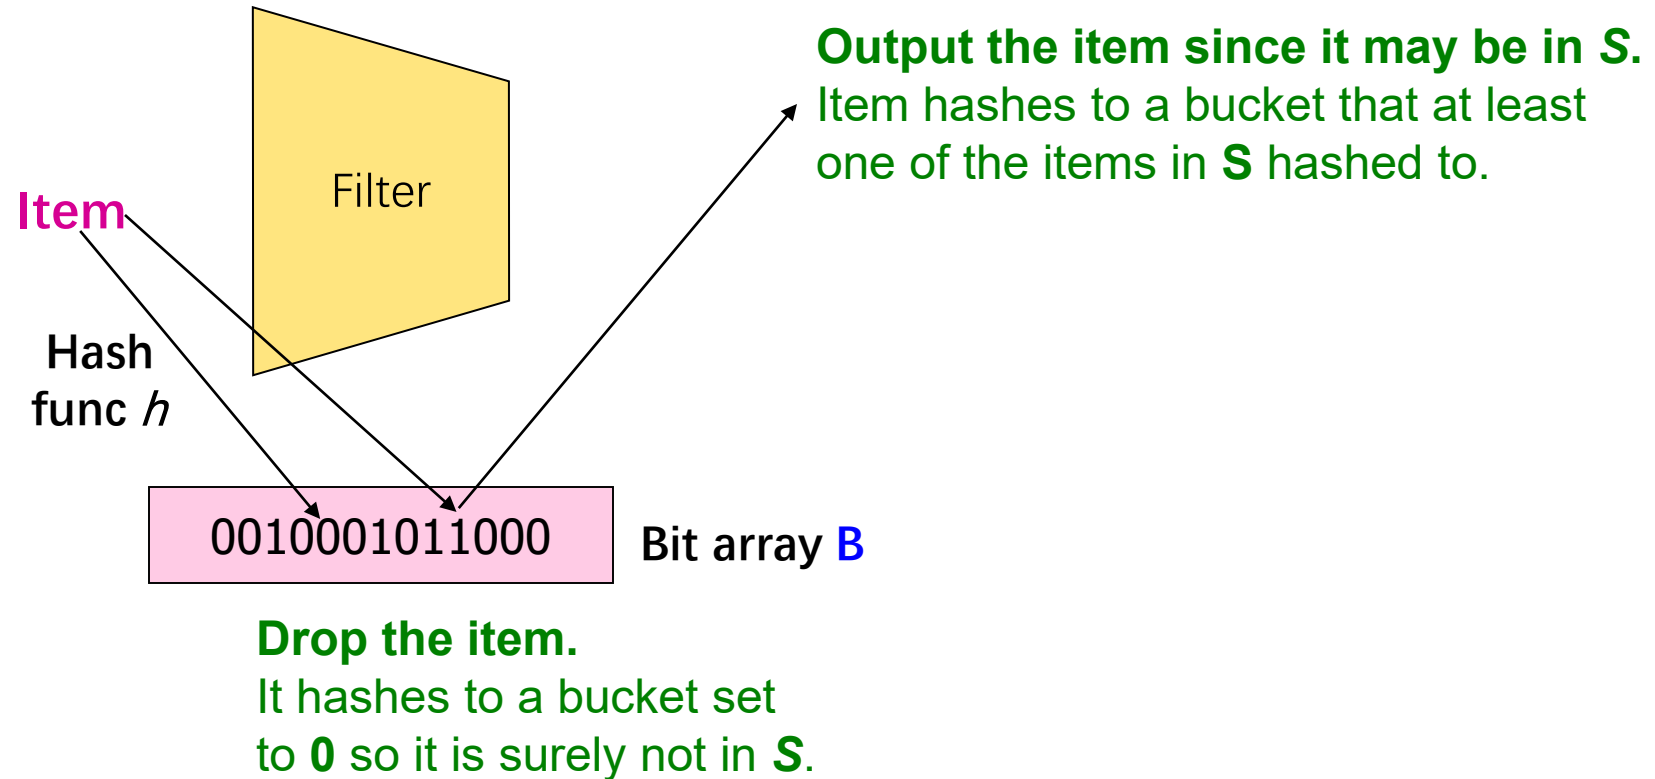
**What can we do?**

# First Cut Solution

- Given a set of keys  $S$  that we want to filter
- Create a **bit array**  $B$  of  $n$  bits, initially all  $0$ s
- Choose a **hash function**  $h$  with range  $[0, n)$
- Hash each member of  $s \in S$  to one of  $n$  buckets, and set that bit to  $1$ , i.e.,  $B[h(s)] = 1$
- Hash each element  $a$  of the stream and output only those that hash to bit that was set to  $1$ 
  - Output  $a$  if  $B[h(a)] = 1$



# First Cut Solution



- Creates **false positives** but **no false negatives**
  - If the item is in  $S$  we surely output it, if not we may still output it

# First Cut Solution

- $|S| = 1$  billion email addresses  
 $|B| = 1\text{GB} = 8$  billion bits, for the hash values
- If the email address is in  $S$ , then it surely hashes to a bucket that has the bit set to **1**, so it always gets through (*no false negatives*)
- Approximately **1/8** of the bits are set to **1**, so about **1/8** of the addresses not in  $S$  get through to the output (*false positives*)


# Analysis: Throwing Darts

- More accurate analysis for the number of **false positives**
- Consider: If we throw  $m$  darts into  $n$  equally likely targets, what is the probability that a target gets at least one dart?
- **In our case:**
  - **Targets** = bits/buckets
  - **Darts** = hash values of items



# Analysis: Throwing Darts

- We have  $m$  darts,  $n$  targets
- What is the probability that **a target gets at least one dart**?




# Analysis: Throwing Darts

- Fraction of 1s in the array B  
= probability of false positive =  $1 - e^{-m/n}$
- **Example:**  $10^9$  darts,  $8 \cdot 10^9$  targets
  - Fraction of **1s** in **B** =  $1 - e^{-1/8} = 0.1175$ 
    - Compare with our earlier estimate:  $1/8 = 0.125$
- How to further **improve** this false positive probability?
- Similar to LSH: Bloom Filter.

# Bloom Filter

- Consider:  $|\mathbf{S}| = m$ ,  $|\mathbf{B}| = n$
- Use  $k$  independent hash functions  $h_1, \dots, h_k$
- **Initialization:**
  - Set  $\mathbf{B}$  to all  $0$ s
  - Hash each element  $s \in \mathcal{S}$  using each hash function  $h_i$ , set  $\mathbf{B}[h_i(s)] = 1$  (for each  $i = 1, \dots, k$ )
- **Run-time:**
  - When a stream element with key  $x$  arrives
    - If  $\mathbf{B}[h_i(x)] = 1$  for all  $i = 1, \dots, k$  then declare that  $x$  is in  $\mathcal{S}$ 
      - That is,  $x$  hashes to a bucket set to  $1$  for every hash function  $h_i(x)$
    - Otherwise discard the element  $x$



**What is the false positive probability?**

# Bloom Filter — Analysis

- **What fraction of the bit vector  $B$  are 1s?**
  - Throwing  $k \cdot m$  darts at  $n$  targets
  - So fraction of 1s is  $(1 - e^{-km/n})$  (false positive of 1 hash function)
- But we have  $k$  independent hash functions and we only let the element  $x$  through **if all**  $k$  hash element  $x$  to a bucket of value **1**
- So, **false positive probability** =  $(1 - e^{-km/n})^k$


# Bloom Filter – Analysis

- $m = 1$  billion,  $n = 8$  billion

- $k = 1$ :  $(1 - e^{-1/8}) = 0.1175$
- $k = 2$ :  $(1 - e^{-1/4})^2 = 0.0493$

- What happens as we keep increasing  $k$ ?


- “Optimal” value of  $k$ :  $n/m \ln(2)$ 
  - In our case: Optimal  $k = 8 \ln(2) = 5.54 \approx 6$ 
    - Error at  $k = 6$ :  $(1 - e^{-1/6})^2 = 0.0235$





# Bloom Filter: Wrap-up

- Bloom filters guarantee **no false negatives**, and use limited memory
  - Great for **pre-processing** before more expensive checks
- Suitable for **hardware** implementation
  - Hash function computations can be **parallelized**
- Disadvantage: only insertion, no deletion from Bloom Filter.




# Count-Min Sketch

# Count Element Frequency

- Faced with big data streams, storing all elements and corresponding frequencies is **impossible**.
- **Approximate** counts are acceptable.
- We can use **hashing** again.


# Approximate Counts with Hashing

- **Initialization**:  $count[i] = 0$ , for  $i \in [1, w]$
- **Increment** count of element  $a$ :  $count[h(a)] += 1$
- **Retrieve** count of element  $a$ :  $count[h(a)]$



# Improvement: More Hash Functions

- We use  $d$  pairwise independent hash functions
- **Increment** count of element  $a$ :  $count[i, h_i(a)] += 1$  for  $i \in [1, d]$
- **Retrieve** count of element  $a$ :  $\min_{i \in [1, d]} count[i, h_i(a)]$



# Guarantees


- Theorem[1]: with probability  $1 - \delta$ , the error is at most  $\varepsilon * \text{count}$ . Concrete values for these error bounds can be chosen by setting  $w = \left\lceil \frac{e}{\varepsilon} \right\rceil$  and  $d = \left\lceil \ln\left(\frac{1}{\delta}\right) \right\rceil$ ,  $e \approx 2.718$ .
  - Adding another **hash** function **exponentially** decreases the chance of hash conflicts
  - Increasing the **width** helps spread up the counts with a **linear** effect

# Queries over a Sliding Window

Streaming Binary Counting

# Sliding Windows

- A useful model of stream processing is that queries are **within** a **window** of length  $N$  – the  $N$  most recent elements received
  - **Amazon example:** For every product  $X$  we keep 0/1 stream of whether that product was **sold** in the  $n$ -th **transaction**. We want answer queries, **how many times we sold  $X$  in the last  $k$  sales.**



Suppose we keep a window with length  $N=6$ ,  
we can query on the last  $k$  transactions, for  $k \leq N$ .



# Counting Bits


- **Problem:**

- Given a stream of **0**s and **1**s
- Be prepared to answer queries of the form  
**How many 1s are in the last  $k$  bits?** where  $k \leq N$

- **Obvious solution:**

Store the most recent  $N$  bits

- When new bit comes in, discard the  $N+1^{\text{st}}$  bit
- **Not feasible** when  $N$  is so **large** that the data cannot be stored in memory, or even on disk




# Counting Bits

- **Real Problem:**

What if we cannot afford to store or compute  $N$  bits?


- **E.g.**, we're processing 1 billion streams and  $N = 1$  billion



- But we are happy with an **approximate** answer

# An attempt: Simple solution

- **Q: How many 1s are in the last  $N$  bits?**
- A simple solution that does **not really** solve our problem:  
**Uniformity assumption**



- **Maintain 2 counters:**
  - $S$ : number of **1**s from the beginning of the stream
  - $Z$ : number of **0**s from the beginning of the stream
- How many 1s are in the last  $N$  bits?  $N \cdot \frac{S}{S+Z}$
- But, what if stream is **non-uniform**?
  - What if distribution changes over time? This is always true in reality.

# DGIM Method

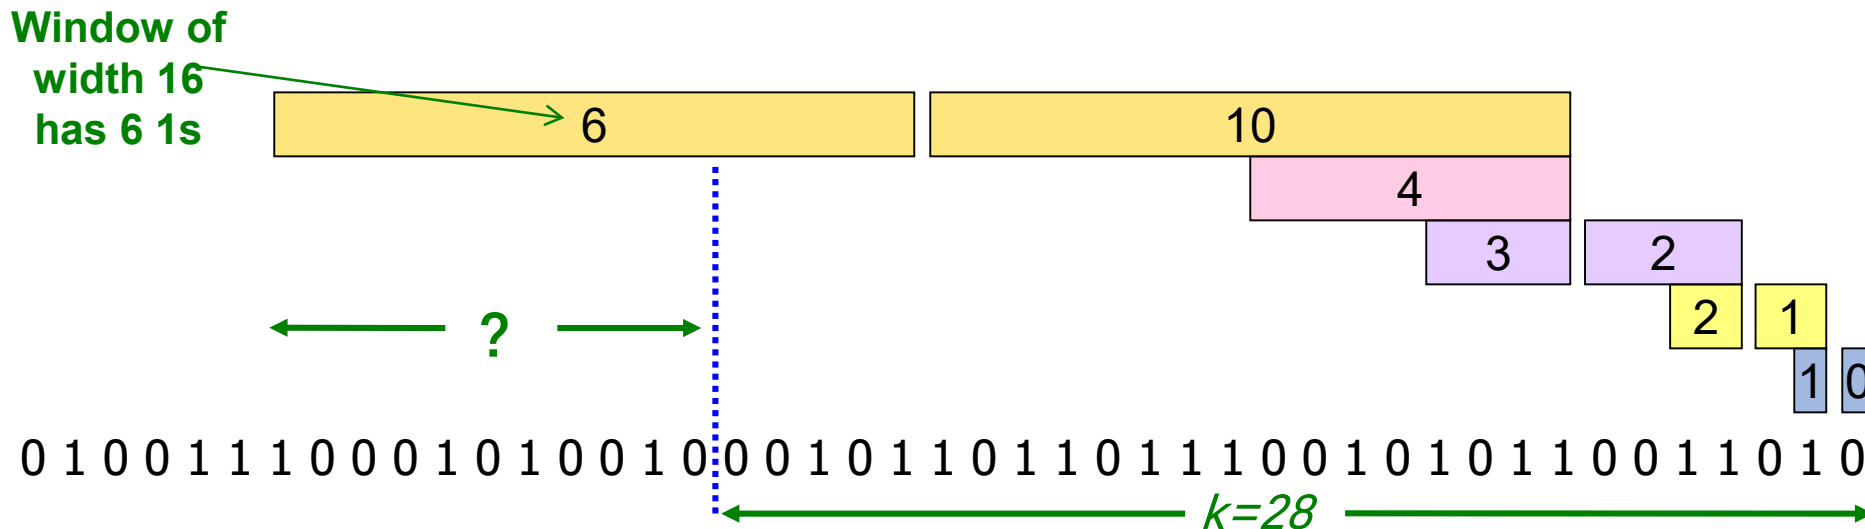
- **DGIM**(*Datar-Gionis-Indyk-Motwani Algorithm*) solution that **does not** assume uniformity
- We store  $\mathbf{O}(\log^2 N)$  bits per stream
- Solution gives **approximate** answer, **never off** by more than **50%**
  - Error factor can be reduced to any fraction  $> 0$ , with more complicated algorithm and proportionally more stored bits

# Idea: Exponential Windows

- **First trial:**

- Summarize **exponentially increasing** regions of the stream, looking backward, to answer queries over last  $k$  items ( $k \leq N$ ).
- Drop small regions if there are more than two on the same level (keep the rightmost)

1. when a bit comes in, create a bucket of length 1 with the proper count (0 or 1).
2. If any level has 3 buckets:
  - a) add the rightmost two and create a bucket at the next higher level (twice the length) with that sum.
  - b) delete the leftmost two buckets, keeping only the rightmost of the three.
3. Repeat (2) recursively for progressively higher levels.



We can reconstruct the count of the last  $k$  bits, except we are not sure how many of the last **6 1s** are included in the  $k=28$  window

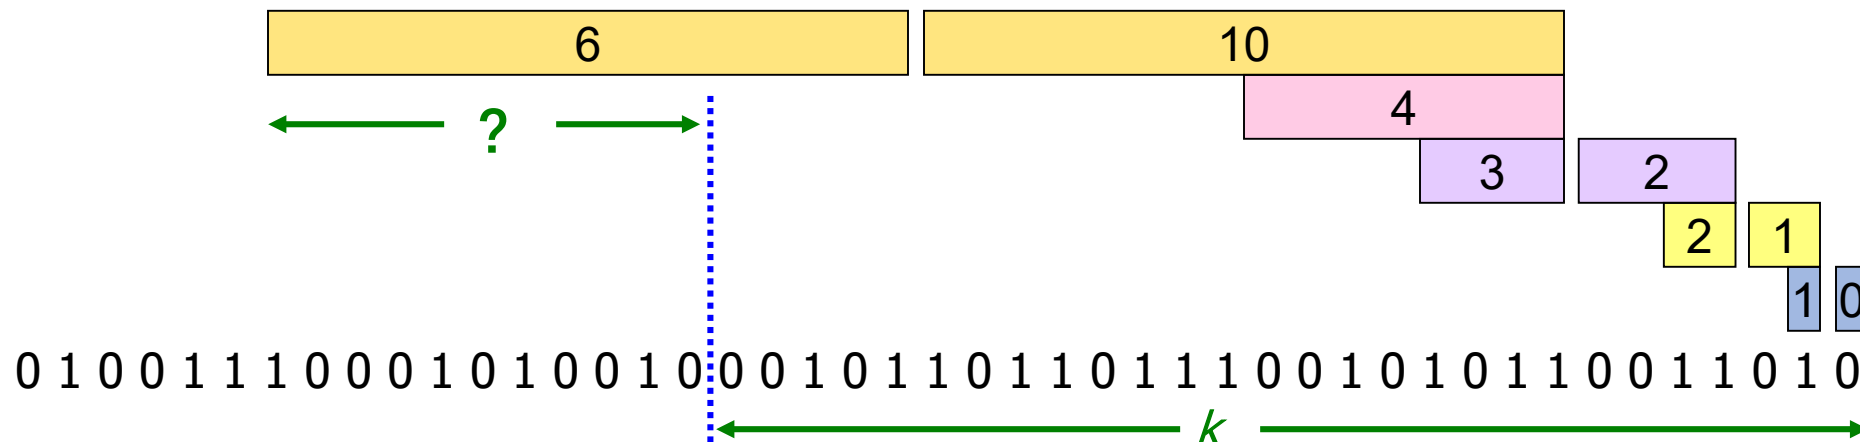
# What's Good?

- Stores only  $O(\log^2 N)$  bits
  - $O(\log N)$  counts of  $\log_2 N$  bits each
- Easy update as more bits enter
- Error in count no greater than the number of **1s** in the “unknown” area

# What's Not So Good?


Relative error = error / true count

- As long as the **1s** are fairly **evenly distributed**, the error due to the unknown region is small
- **What is the bound of the relative error?**
  - Consider the case that all the **1s** are in the unknown area(?) part) and the rest are all 0s. Here the relative error is infinite.



# Fixup: DGIM method

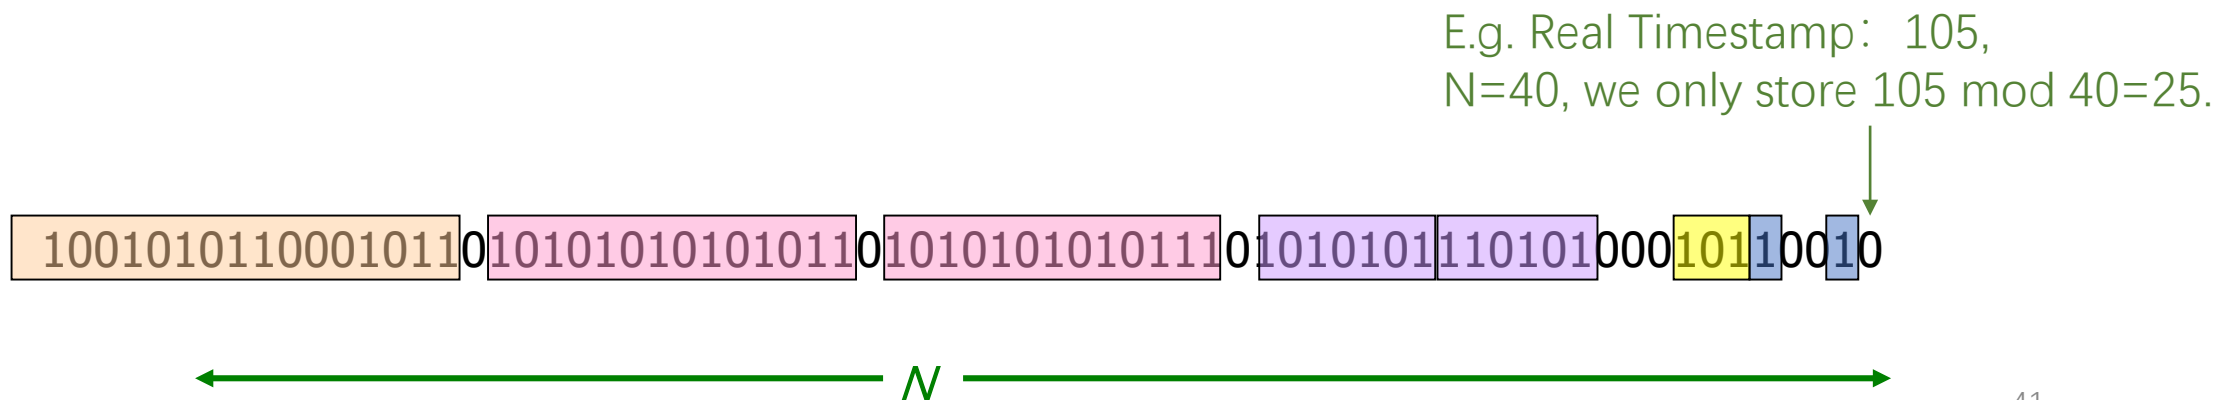
- **Idea:** Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
  - Let the block **sizes** (number of **1s**) increase **exponentially**
  - **Data dependent**
- When there are few 1s in the window, block sizes stay small, so errors are small





# DGIM: Timestamps

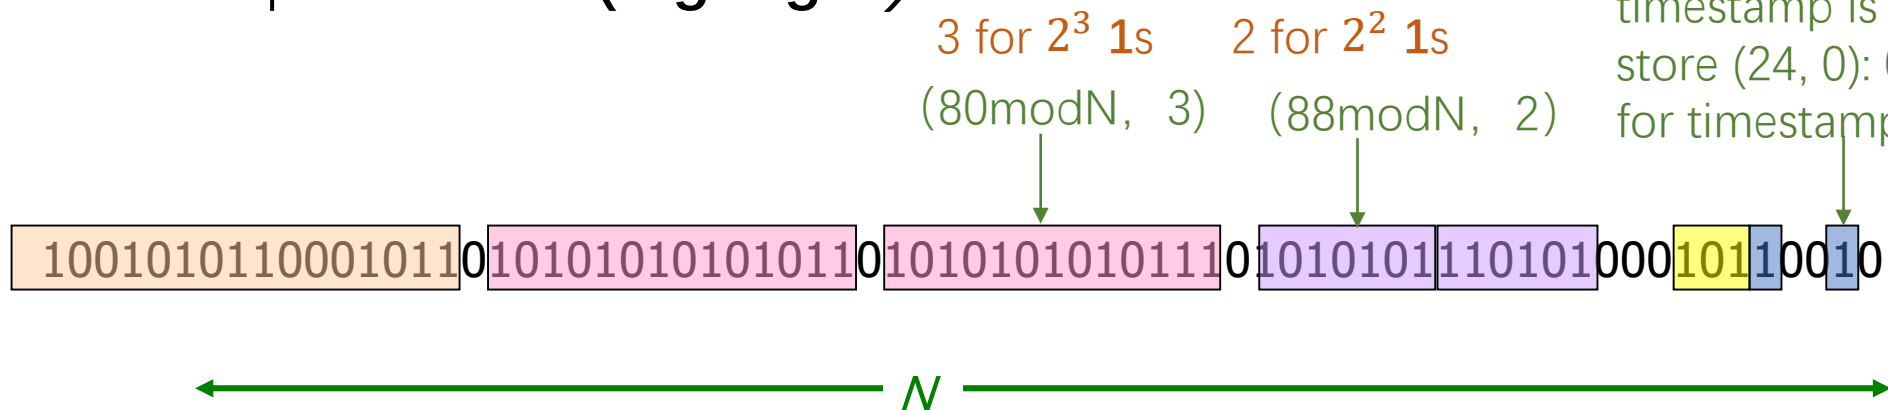
- Each bit in the stream has a *timestamp*, starting **1, 2, ...**
- Record timestamps **modulo  $N$**  (**the window size**), so we can represent any **relevant** timestamp in  **$O(\log_2 N)$**  bits



# DGIM: Buckets

- A *bucket* in the DGIM method is a record consisting of:
  - The timestamp of its end [ $O(\log N)$  bits]
  - The number of 1s between its beginning and end [ $O(\log \log N)$  bits]


- **Constraint on buckets:**  
Number of **1s** must be a power of 2
  - That explains the  $O(\log \log N)$



E.g. In this window, if the timestamp of the last timestamp is 105, we actually store (24, 0): 0 for  $2^0 = 1$ , 24 for timestamp 104 mod 40.

# Representing a Stream by Buckets

- Either **one** or **two** buckets with the same **power-of-2 number** of **1s**
- Buckets **do not overlap** in timestamps
- Buckets are **sorted** by **size**
  - Earlier buckets are not smaller than later buckets
- Buckets **disappear** when their end-time is **>  $N$**  time units in the past



# Updating Buckets

- When a new bit comes in, **drop** the last (oldest) bucket if its end-time is **prior to  $N$**  time units before the current time
- **2 cases:** Current bit is **0** or **1**
- **If the current bit is 0:**  
**no other changes are needed**

# Updating Buckets

- **If the current bit is 1:**
  - **(1)** Create a new bucket of size **1**, for just this bit
    - **End timestamp = current time**
  - **(2)** If there are now three buckets of size 1, combine the oldest two into a bucket of size 2
  - **(3)** If there are now three buckets of size 2, combine the oldest two into a bucket of size 4
  - **(4)** And so on ...

# Example: Updating Buckets

**Current state of the stream:**

100101011000101101010101010101101010101010111010101011101010111010100010110010

**Bit of value 1 arrives**

0010101100010110101010101010110101010101110101010111010101110101000101100101

**Two smallest buckets get merged into a size-2 bucket**

001010110001011010101010101011010101010111010101110101000101100101

**Next bit 1 arrives, new size-1 bucket is created, then 0 comes, then 1:**

010110001011010101010101011010101010111010101110101000101100101101

**Buckets get merged...**


010110001011010101010101011010101010111010101110101000101100101101

**State of the buckets after merging**

010110001011010101010101011010101010111010101110101000101100101101


# How to Query?

- To estimate the number of 1s in the most recent  $k$  bits:
  1. Sum the sizes of **all** buckets **but the last**
  2. Add **half** the size of the last bucket
- **Remember:** We do **not** know how many **1s** of the last bucket are still within the wanted window



# Error Bound: Proof


- **Why is error 50%? Let's prove it!**
- Suppose the last bucket has size  $2^r$
- Then by assuming  $2^{r-1}$  (i.e., half) of its **1s** are still within the window, we make an error of at most  $2^{r-1}$
- Since there is at least one bucket of each of the sizes less than  $2^r$ , the true sum is at least  
 $1 + 2 + 4 + \dots + 2^{r-1} = 2^r - 1$
- Thus, **relative error at most 50%**





# Further Reducing the Error

- Instead of maintaining **1** or **2** of each size bucket, **we allow either  $s-1$  or  $s$  buckets ( $s > 2$ )**
  - Except for the largest size buckets, where we can have any number between **1** and  **$s$**  of those.



- **Error is at most  $\frac{2^{r-1}}{(s-1)(2^r-1)} = O(1/s)$**
- By picking  **$s$**  appropriately, we can tradeoff between number of bits we store and the error

# Extensions

- Can we handle the case where the stream is not bits, but **integers**, and we want the **sum** of the last  $k$  elements?
- We want the sum of the last  $k$  elements
  - Amazon: Avg. price of last  $k$  sales
- **Solution:**
  - If you know all have at most  $m$  bits
    - Treat  $m$  bits of each integer as a separate stream
    - Use DGIM to count **1s** in each integer
    - The sum is  $= \sum_{i=0}^{m-1} c_i 2^i$

11111111000000001110101010101101010101010111010101011101010001011001**0**  
11111111000000001110101010101101010100010111010101011101010001111001**1**

Two streams represent 1